

# Machine Learning

## Part 1: Mathematical Foundation of Machine Learning

---

Zengchang Qin (PhD)

# Probability & Statistics

---

# Probability (Objective and Subjective)


---

The first approach is to define probability in terms of frequency of occurrence, as a percentage of successes in a moderately large number of similar situations.



Such an interpretation is often natural. For example, when we say that a perfectly manufactured coin lands on heads “with probability 50%,” we typically mean “roughly half of the time.”

Consider, for example, a scholar who asserts that the Iliad and the Odyssey were composed by the same person, with probability 90%. Such an assertion conveys some information, but not in terms of frequencies, since the subject is a one-time event. Rather, it is an expression of the scholar’s subjective belief.




# Set Operation


---

## Examples of Venn diagrams.


- (a) The shaded region is  $S \cap T$ .
- (b) The shaded region is  $S \cup T$ .
- (c) The shaded region is  $S \cap c(T)$ .
- (d) Here,  $T \subset S$ . The shaded region is the complement of  $S$ .
- (e) The sets  $S$ ,  $T$ , and  $U$  are disjoint.
- (f) The sets  $S$ ,  $T$ , and  $U$  form a partition of the set  $\Omega$




(a)




(b)




(c)



(d)



(e)



(f)

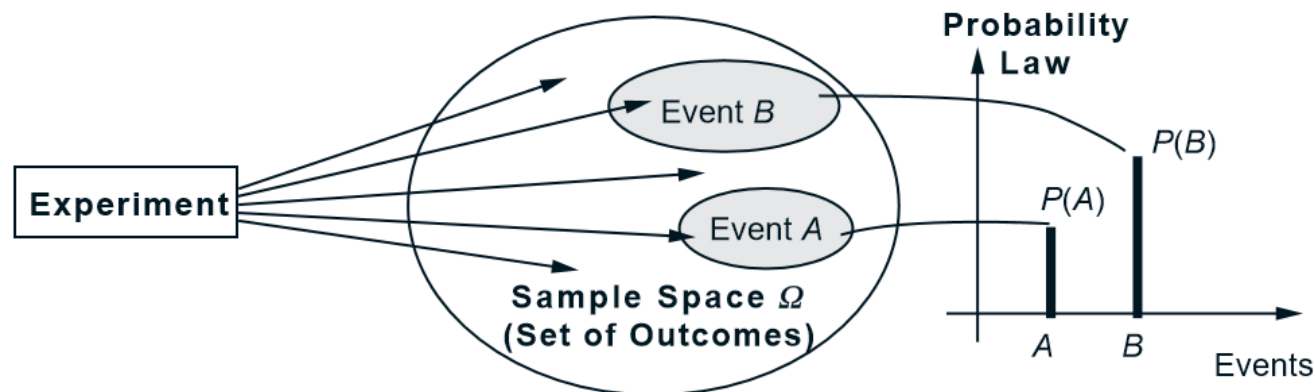
# Probabilistic Models

---

## Elements of a Probabilistic Model

The sample space  $\Omega$ , which is the set of all possible outcomes of an experiment.

The **probability law**, which assigns to a set  $A$  of possible outcomes (also called an event) a nonnegative number  $P(A)$  (called the probability of  $A$ ) that encodes our knowledge or belief about the collective “likelihood” of the elements of  $A$ . The probability law must satisfy certain properties to be introduced shortly.



# Probability Axioms

---

## Probability Axioms

1. (**Nonnegativity**)  $\mathbf{P}(A) \geq 0$ , for every event  $A$ .
2. (**Additivity**) If  $A$  and  $B$  are two disjoint events, then the probability of their union satisfies

$$\mathbf{P}(A \cup B) = \mathbf{P}(A) + \mathbf{P}(B).$$

Furthermore, if the sample space has an infinite number of elements and  $A_1, A_2, \dots$  is a sequence of disjoint events, then the probability of their union satisfies

$$\mathbf{P}(A_1 \cup A_2 \cup \dots) = \mathbf{P}(A_1) + \mathbf{P}(A_2) + \dots$$

3. (**Normalization**) The probability of the entire sample space  $\Omega$  is equal to 1, that is,  $\mathbf{P}(\Omega) = 1$ .

# Conditional Probability

---

## Properties of Conditional Probability

- The conditional probability of an event  $A$ , given an event  $B$  with  $\mathbf{P}(B) > 0$ , is defined by

$$\mathbf{P}(A | B) = \frac{\mathbf{P}(A \cap B)}{\mathbf{P}(B)},$$

and specifies a new (conditional) probability law on the same sample space  $\Omega$ . In particular, all known properties of probability laws remain valid for conditional probability laws.

- Conditional probabilities can also be viewed as a probability law on a new universe  $B$ , because all of the conditional probability is concentrated on  $B$ .
- In the case where the possible outcomes are finitely many and equally likely, we have

$$\mathbf{P}(A | B) = \frac{\text{number of elements of } A \cap B}{\text{number of elements of } B}.$$



Let's consider a problem of **conditional probability**:

My neighbor John has two kids.

1. He told me that one of his two kids is a boy, what is the probability that the other one is a girl.
2. If I saw one's kids is playing outside, that is a boy, what is the probability that the other one is a girl.

# Total Probability Theorem

---



## Total Probability Theorem

Let  $A_1, \dots, A_n$  be disjoint events that form a partition of the sample space (each possible outcome is included in one and only one of the events  $A_1, \dots, A_n$ ) and assume that  $\mathbf{P}(A_i) > 0$ , for all  $i = 1, \dots, n$ . Then, for any event  $B$ , we have

$$\begin{aligned}\mathbf{P}(B) &= \mathbf{P}(A_1 \cap B) + \dots + \mathbf{P}(A_n \cap B) \\ &= \mathbf{P}(A_1)\mathbf{P}(B | A_1) + \dots + \mathbf{P}(A_n)\mathbf{P}(B | A_n).\end{aligned}$$



# Independence

---

## Bayes' Rule

Let  $A_1, A_2, \dots, A_n$  be disjoint events that form a partition of space, and assume that  $\mathbf{P}(A_i) > 0$ , for all  $i$ . Then, for any event  $B$  that  $\mathbf{P}(B) > 0$ , we have

$$\begin{aligned}\mathbf{P}(A_i | B) &= \frac{\mathbf{P}(A_i)\mathbf{P}(B | A_i)}{\mathbf{P}(B)} \\ &= \frac{\mathbf{P}(A_i)\mathbf{P}(B | A_i)}{\mathbf{P}(A_1)\mathbf{P}(B | A_1) + \dots + \mathbf{P}(A_n)\mathbf{P}(B | A_n)}\end{aligned}$$



REV. T. BAYES

## Independence

- Two events  $A$  and  $B$  are said to be independent if

$$\mathbf{P}(A \cap B) = \mathbf{P}(A)\mathbf{P}(B).$$

If in addition,  $\mathbf{P}(B) > 0$ , independence is equivalent to the condition

$$\mathbf{P}(A | B) = \mathbf{P}(A).$$

- If  $A$  and  $B$  are independent, so are  $A$  and  $B^c$ .
- Two events  $A$  and  $B$  are said to be conditionally independent, given another event  $C$  with  $\mathbf{P}(C) > 0$ , if

$$\mathbf{P}(A \cap B | C) = \mathbf{P}(A | C)\mathbf{P}(B | C).$$


If in addition,  $\mathbf{P}(B \cap C) > 0$ , conditional independence is equivalent to the condition

$$\mathbf{P}(A | B \cap C) = \mathbf{P}(A | C).$$


- Independence does not imply conditional independence, and vice versa.

# Random Variable

---



(a)



(b)

(a) Visualization of a random variable. It is a function that assigns a numerical value to each possible outcome of the experiment.


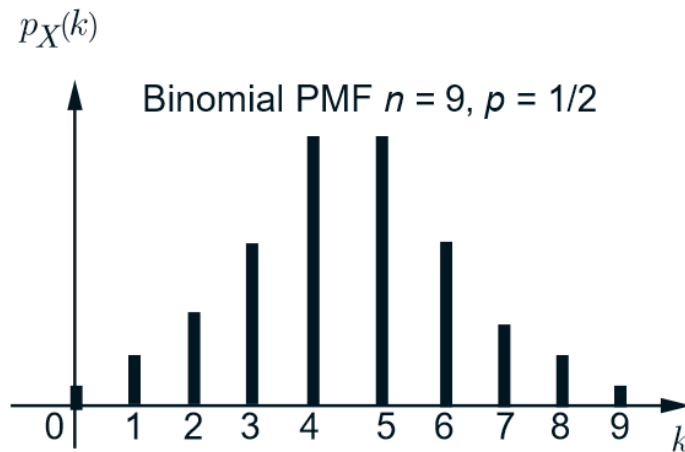
(b) An example of a random variable. The experiment consists of two rolls of a 4-sided die, and the random variable is the maximum of the two rolls. If the outcome of the experiment is (4,2), the experimental value of this random variable is 4.

# Binomial Random Variable

---

At each toss, the coin comes up a head with probability  $p$ , and a tail with probability  $1-p$ , independently of prior tosses. Let  $X$  be the number of heads in the  $n$ -toss sequence. We refer to  $X$  as a binomial random variable with parameters  $n$  and  $p$ .

$$p_X(k) = \mathbf{P}(X = k) = \binom{n}{k} p^k (1-p)^{n-k}, \quad k = 0, 1, \dots, n. \quad \sum_{k=0}^n \binom{n}{k} p^k (1-p)^{n-k} = 1.$$




# Binomial Probabilities

---

We showed above that the probability of any given sequence that contains  $k$  heads is  $p^k(1-p)^{n-k}$ , so we have

$$p(k) = \binom{n}{k} p^k (1-p)^{n-k},$$

where  $\binom{n}{k}$  = number of distinct  $n$ -toss sequences that contain  $k$  heads.



The numbers  $\binom{n}{k}$  (called “ $n$  choose  $k$ ”) are known as the **binomial coefficients**, while the probabilities  $p(k)$  are known as the **binomial probabilities**. Using a counting argument, to be given in Section 1.6, one finds that

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}, \quad k = 0, 1, \dots, n,$$

# Expectation of Random Variable and Function of Random Variables

---

## Expectation

We define the **expected value** (also called the **expectation** or the **mean**) of a random variable  $X$ , with PMF  $p_X(x)$ , by<sup>†</sup>

$$\mathbf{E}[X] = \sum_x xp_X(x).$$

## Expected Value Rule for Functions of Random Variables

Let  $X$  be a random variable with PMF  $p_X(x)$ , and let  $g(X)$  be a real-valued function of  $X$ . Then, the expected value of the random variable  $g(X)$  is given by

$$\mathbf{E}[g(X)] = \sum_x g(x)p_X(x).$$

# Variance

---

## Variance

The variance  $\text{var}(X)$  of a random variable  $X$  is defined by

$$\text{var}(X) = \mathbf{E}[(X - \mathbf{E}[X])^2]$$

and can be calculated as

$$\text{var}(X) = \sum_x (x - \mathbf{E}[X])^2 p_X(x).$$

It is always nonnegative. Its square root is denoted by  $\sigma_X$  and is called the **standard deviation**.

# Covariance

---




Covariance is a measure of the linear relationship between two random variables. We denote the covariance between  $X$  and  $Y$  as  $\text{Cov}(X, Y)$ , and it is defined to be

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

Note that the outer expectation must be taken over the joint distribution of  $X$  and  $Y$ .

Again, the linearity of expectation allows us to rewrite this as

$$\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$$



# Joint Probability of More than One Random Variables

---

Joint PMF  $P_{X,Y}(x,y)$   
in tabular form

4	0	1/20	1/20	1/20	3/20
3	1/20	2/20	3/20	1/20	7/20
2	1/20	2/20	3/20	1/20	7/20
1	1/20	1/20	1/20	0	3/20
	1	2	3	4	
	3/20	6/20	8/20	3/20	

Column Sums:  
Marginal PMF  $P_X(x)$

$$p_{X,Y,Z}(x, y, z) = \mathbf{P}(X = x, Y = y, Z = z),$$

$$p_X(x) = \sum_y \sum_z p_{X,Y,Z}(x, y, z).$$

Row Sums:  
Marginal PMF  $P_Y(y)$

$$p_{X,Y}(x, y) = \sum_z p_{X,Y,Z}(x, y, z),$$


$$\mathbf{E}[g(X, Y, Z)] = \sum_{x,y,z} g(x, y, z) p_{X,Y,Z}(x, y, z),$$



# Conditioning

The world is full of conditions, when we say independent, it usually implies “conditional independent”.

Independent or dependent?



Visualization of the conditional PMF  $p_{X|Y}(x|y)$ . For each  $y$ , we view the joint PMF along the slice  $Y = y$  and renormalize so that

$$\sum_x p_{X|Y}(x|y) = 1.$$

## Continuous PDF

---

a **probability distribution** is a mathematical function that, stated in simple terms, can be thought of as providing the probabilities of occurrence of different possible outcomes in an experiment. From frequency to a continuous function.






Illustration of a PDF. The probability that  $X$  takes value in an interval  $[a, b]$  is  $\int_a^b f_X(x) dx$ , which is the shaded area in the figure.

# Probability Distributions

---



Interpretation of the PDF  $f_X(x)$  as “probability mass per unit length” around  $x$ . If  $\delta$  is very small, the probability that  $X$  takes value in the interval  $[x, x + \delta]$  is the shaded area in the figure, which is approximately equal to  $f_X(x) \cdot \delta$ .



# Bernoulli Distribution

---

The probability distribution of any single experiment that asks a yes–no question; the question results in a Boolean valued outcome, a single bit of information whose value is success/yes/true/one with probability  $p$  and failure/no/false/zero with probability  $q$ .

If  $X$  is a random variable with this distribution, we have:

$$\Pr(X = 1) = p = 1 - \Pr(X = 0) = 1 - q.$$

The **probability mass function**  $f$  of this distribution, over possible outcomes  $k$ , is

$$f(k; p) = \begin{cases} p & \text{if } k = 1, \\ 1 - p & \text{if } k = 0. \end{cases}$$

This can also be expressed as

$$f(k; p) = p^k (1 - p)^{1-k} \quad \text{for } k \in \{0, 1\}$$

$$\mathbb{E}(X) = p \quad \mathbb{E}[X^2] = \Pr(X = 1) \cdot 1^2 + \Pr(X = 0) \cdot 0^2 = p \cdot 1^2 + q \cdot 0^2 = p$$

$$\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = p - p^2 = p(1 - p) = pq$$

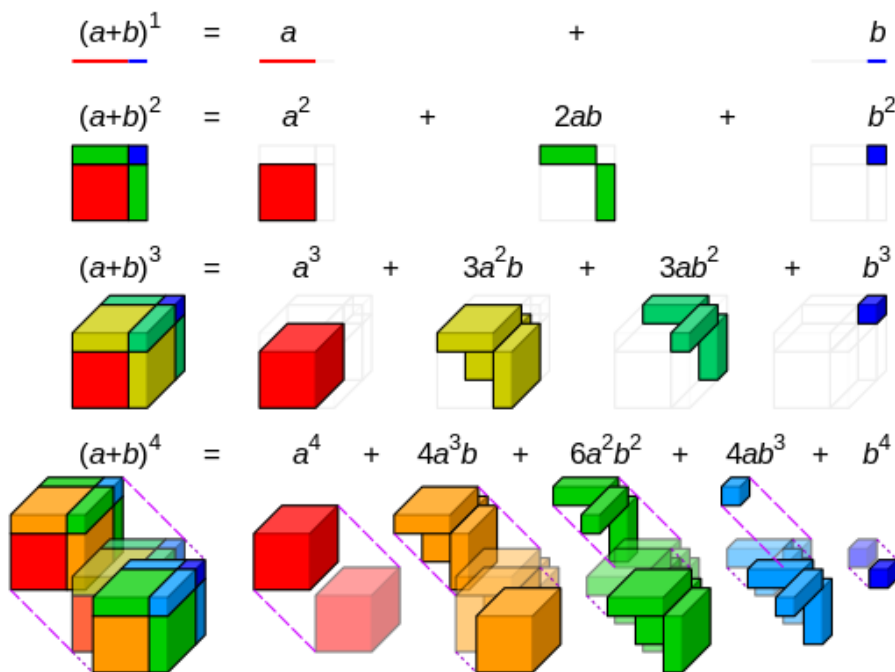


Jacob Bernoulli

# Binomial Distribution

---

The binomial distribution with parameters  $n$  and  $p$  is the discrete probability distribution of the number of successes in a sequence of  $n$  independent Bernoulli trials of yes–no questions.

$$\begin{aligned}
 (a+b)^1 &= a + b \\
 (a+b)^2 &= a^2 + 2ab + b^2 \\
 (a+b)^3 &= a^3 + 3a^2b + 3ab^2 + b^3 \\
 (a+b)^4 &= a^4 + 4a^3b + 6a^2b^2 + 4ab^3 + b^4
 \end{aligned}$$


$$(x+y)^n = \sum_{k=0}^n \binom{n}{k} x^{n-k} y^k = \sum_{k=0}^n \binom{n}{k} x^k y^{n-k}$$

$$\begin{array}{cccccccc}
 & & & & 1 & & & \\
 & & & 1 & & 1 & & \\
 & & 1 & & 2 & & 1 & \\
 & 1 & & 3 & & 3 & & 1 \\
 1 & & 1 & & 4 & & 6 & & 4 & & 1 \\
 & 1 & & 5 & & 10 & & 10 & & 5 & & 1 \\
 & & 1 & & 6 & & 15 & & 20 & & 15 & & 6 & & 1 \\
 1 & & & 1 & & 7 & & 21 & & 35 & & 35 & & 21 & & 7 & & 1
 \end{array}$$

Pascal's triangle

# Multinomial Distribution

---

$$f(x_1, \dots, x_k; n, p_1, \dots, p_k) = \Pr(X_1 = x_1 \text{ and } \dots \text{ and } X_k = x_k) \\ = \begin{cases} \frac{n!}{x_1! \dots x_k!} p_1^{x_1} \times \dots \times p_k^{x_k}, & \text{when } \sum_{i=1}^k x_i = n \\ 0 & \text{otherwise,} \end{cases}$$

Suppose one does an experiment of extracting  $n$  balls of  $k$  different colours from a bag, replacing the extracted ball after each draw. Balls from the same colour are equivalent.

The probability mass function can be expressed using the gamma function as:

$$f(x_1, \dots, x_k; p_1, \dots, p_k) = \frac{\Gamma(\sum_i x_i + 1)}{\prod_i \Gamma(x_i + 1)} \prod_{i=1}^k p_i^{x_i}$$


Where gamma function is an extension of factorial:

$$\Gamma(n) = (n-1)!$$

$$\Gamma(z) = \int_0^\infty x^{z-1} e^{-x} dx$$

## More on Gamma Function

It is easy graphically to interpolate the factorial function to non-integer values, but is there a formula that describes the resulting curve?



$$\begin{aligned}\Gamma(z+1) &= \int_0^{\infty} x^z e^{-x} dx \\ &= [-x^z e^{-x}]_0^{\infty} + \int_0^{\infty} z x^{z-1} e^{-x} dx \\ &= \lim_{x \rightarrow \infty} (-x^z e^{-x}) - (0e^{-0}) + z \int_0^{\infty} x^{z-1} e^{-x} dx\end{aligned}$$

Recognizing that as  $x \rightarrow \infty$ ,  $-x^z e^{-x} \rightarrow 0$ ,

$$\Gamma(z+1) = z \int_0^{\infty} x^{z-1} e^{-x} dx = z\Gamma(z)$$

$$\begin{aligned}\Gamma(1) &= \int_0^{\infty} x^{1-1} e^{-x} dx = [-e^{-x}]_0^{\infty} \\ &= \lim_{x \rightarrow \infty} (-e^{-x}) - (-e^{-0}) = 0 - (-1) = 1\end{aligned}$$

Given that  $\Gamma(1) = 1$  and  $\Gamma(n+1) = n\Gamma(n)$

$$\Gamma(n) = 1 \cdot 2 \cdot 3 \cdots (n-1) = (n-1)!$$

# Gamma Distribution

---


Like the lognormal the gamma distribution is unbounded on the right, defined for only positive  $X$ , and tends to yield skewed distributions.

$$X \sim \Gamma(\alpha, \beta) \equiv \text{Gamma}(\alpha, \beta)$$

The corresponding **probability density function** in the shape-rate parametrization is

$$f(x; \alpha, \beta) = \frac{\beta^\alpha x^{\alpha-1} e^{-\beta x}}{\Gamma(\alpha)} \quad \text{for } x > 0 \text{ and } \alpha, \beta > 0,$$

where  $\Gamma(\alpha)$  is a complete **gamma function**.



The gamma distribution is widely used as a conjugate prior in Bayesian statistics. It is the conjugate prior for the precision of a normal distribution. It is also the conjugate prior for the exponential distribution.




# Beta Distribution

---

It is bounded on both sides. In this respect it resembles the binomial distribution. The standard beta distribution is constrained so that its domain is the interval  $(0, 1)$ .

$$\begin{aligned} f(x; \alpha, \beta) &= \text{constant} \cdot x^{\alpha-1} (1-x)^{\beta-1} \\ &= \frac{x^{\alpha-1} (1-x)^{\beta-1}}{\int_0^1 u^{\alpha-1} (1-u)^{\beta-1} du} = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1} \\ &= \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1} \end{aligned}$$



The beta function, **B** is a normalization constant to ensure that the total probability integrates to 1.

# Poisson Distribution

---

A discrete random variable  $X$  is said to have a Poisson distribution with parameter  $\lambda > 0$ , if, for  $k = 0, 1, 2, \dots$ ,


For example, on a particular river, overflow floods occur once every 100 years on average. Calculate the probability of  $k = 0, 1, 2, 3, 4, 5$ , or 6 overflow floods in a 100-year interval, assuming the Poisson model is appropriate. Because the average event rate is one overflow flood per 100 years,  $\lambda = 1$ , so that:

$$P(k \text{ overflow floods in 100 years}) = \frac{\lambda^k e^{-\lambda}}{k!} = \frac{1^k e^{-1}}{k!}$$

$$P(k = 0 \text{ overflow floods in 100 years}) = \frac{1^0 e^{-1}}{0!} = \frac{e^{-1}}{1} = 0.368$$



$$P(k = 1 \text{ overflow flood in 100 years}) = \frac{1^1 e^{-1}}{1!} = \frac{e^{-1}}{1} = 0.368$$

$$P(k = 2 \text{ overflow floods in 100 years}) = \frac{1^2 e^{-1}}{2!} = \frac{e^{-1}}{2} = 0.184$$



$$f(k; \lambda) = \Pr(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

# Gaussian (Normal) Distribution



The probability density of the normal distribution is:



$$f(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- $\mu$  is the mean or expectation of the distribution
- $\sigma$  is the standard deviation
- $\sigma^2$  is the variance

# Log Normal Distribution

Given a log-normally distributed random variable  $X$  and two parameters  $\mu$  and  $\sigma$  that are, respectively, the **mean** and **standard deviation** of the variable's natural **logarithm**, then the logarithm of  $X$  is normally distributed, and we can write  $X$  as

$$X = e^{\mu + \sigma Z} \quad \text{with } Z \text{ a standard normal variable.}$$





# Exponential Distribution

The exponential distribution (also known as negative exponential distribution) is the probability distribution that describes the time between events in a Poisson process, i.e. a process in which events occur continuously and independently at a constant average rate. It is a particular case of the gamma distribution. It is the continuous analogue of the geometric distribution, and it has the key property of being memoryless.

The **probability density function** (pdf) of an exponential distribution is

$$f(x; \lambda) = \begin{cases} \lambda e^{-\lambda x} & x \geq 0, \\ 0 & x < 0. \end{cases} \quad \mathbb{E}[X] = \frac{1}{\lambda} \quad \text{Var}[X] = \frac{1}{\lambda^2}$$

The exponential distribution occurs naturally when describing the lengths of the inter-arrival times in a homogeneous Poisson process.



## Bayesian Examples

---

Marie is getting married tomorrow at an outdoor ceremony in the desert. In recent years, it has rained only **5 days each year**. Unfortunately, the weatherman is forecasting rain for tomorrow. When it actually rains, the weatherman has forecast rain **90% of the time**. When it doesn't rain, he has forecast rain **10%** of the time. What is the probability it will rain on the day of Marie probability it will rain on the day of Marie s' wedding?

**Event A:** The weatherman has forecast rain.

**Event B:** It rains.

We know:


1.  $p(B) = 5 / 365 = 0.0137$  [ It rains 5 days out of the year. ]
2.  $p(\text{not } B) = 360 / 365 = 0.9863$
3.  $p(A | B) = 0.9$  [ When it rains, the weatherman has forecast rain 90% of the time. ]
4.  $p(A | \text{not } B) = 0.1$  [When it does not rain weatherman has forecast rain 10% of the time.]

# Bayesian Example

---

We want to know  $p(B | A)$ , the probability it will rain on the day of Marie's wedding, given a forecast for rain by the weatherman. The answer can be determined from Bayes Rule:

1.  $p(B | A) = p(A | B) \cdot p(B) / p(A)$
2.  $p(A) = p(A | B) \cdot p(B) + p(A | \text{not } B) \cdot p(\text{not } B)$   
 $= (0.9)(0.014) + (0.1)(0.986) = 0.111$
1.  $p(B | A) = (0.9)(0.0137) / 0.111 = 0.111$



The **special** coin problem.

# Simpson's Paradox

---

**Simpson's paradox**, or the Yule–Simpson effect, is a phenomenon in probability and statistics, in which a trend appears in different groups of data but disappears or reverses when these groups are combined.

Department	Female Applicants	Female Admitted	%	Male Applicants	Male Admitted	%	All Applicants	All Admitted	Overall %
<b>Business School</b>	100	49	49%	20	15	75%	<b>120</b>	<b>64</b>	<b>53.3%</b>
<b>Law School</b>	20	1	5%	100	10	10%	<b>120</b>	<b>11</b>	<b>9.2%</b>
<b>Both</b>	120	50	42%	120	25	21%	<b>240</b>	<b>75</b>	<b>31.3%</b>

Suppose two people, Lisa and Bart, each edit articles for two weeks. In the first week, Lisa fails to improve the only article she edited, and Bart improves 1 of the 4 articles he edited. In the second week, Lisa improves 3 of 4 articles she edited, while Bart improves the only article he edited.

	Week 1	Week 2	Total
<b>Lisa</b>	0/1	3/4	<b>3/5</b>
<b>Bart</b>	<b>1/4</b>	<b>1/1</b>	<b>2/5</b>